




Magazine Article / Corporate Strategy

AI Is Revolutionizing Strategic Decision-Making

New tools can improve human judgment by tirelessly generating, evaluating, and synthesizing insights. *by Felipe A. Csaszar*

From the Magazine (September–October 2026) / Reprint R2605B



Dimitris Ladopoulos

Think about your last strategy offsite. Your team spent weeks preparing. You flew people in, booked the conference room, hired a facilitator. And after two days of debate, how many truly different strategic options did you walk out with? Three? Four? Now ask yourself: Was that because only three or four good options existed, or because that's all your team had the time and mental bandwidth to develop and evaluate?

For most companies, the honest answer is the latter. The bottleneck in strategic decision-making has never been a shortage of possible directions. It has been the limited capacity of the human minds doing the work. We can hold only so much information in our heads, evaluate only so many alternatives in a strategic-planning cycle, and process only so many perspectives in a meeting before fatigue, politics, or the clock forces a decision. Scholars call this dynamic *bounded rationality*—the idea that human decision-makers, however capable, are constrained by finite attention, memory, and processing power.

These constraints are so fundamental that we rarely notice them, but they have quietly shaped every tool in the standard strategy playbook. The reason a SWOT analysis has four quadrants, a growth share matrix is a 2×2, and Michael Porter's most famous framework has exactly five forces is not that the competitive world is actually so simple. It's that the frameworks had to be simple enough for a human team to map them out on a whiteboard in a few hours.

For decades, those frameworks were the best we had. That's no longer the case.

The current generation of artificial intelligence tools—particularly large language models (LLMs) and the multiagent systems being built on top of them—are not just additions to the planning tool kit. They're technologies that directly relax the cognitive constraints that have shaped how companies make their most important decisions. AI can generate and screen thousands of strategic alternatives where a human team might be able to consider a mere handful. It can build and continuously update models of markets, customers, and competitors that are far richer and more dynamic than any static framework. And it can test ideas through simulated deliberation—synthesizing diverse

perspectives and challenging assumptions without the groupthink, hierarchy, and time pressure that distort real-world strategy meetings.

In short, AI offers a path to *unbounding* the strategy process. The limits are not removed, of course, but they're pushed outward—and where they end up matters enormously, because the implications go well beyond efficiency. When you can explore more options, you find better paths forward. When you can model your environment in higher resolution, you see more opportunities and threats. When you can stress-test a plan by simulating competitors, skeptical customers, and a devil's advocate who never gets tired, you make decisions that are more resilient. The companies that master this new way of working won't just do strategy faster. They'll do it *better*—and build the next generation of competitive advantage.

I study strategic decision-making and have spent considerable time examining how AI performs on core strategy tasks. In recent experiments my colleagues and I found that business plans generated by an LLM were rated more favorably by experienced investors than were plans written by entrepreneurs in a startup accelerator and business plan competition. Perhaps more striking, when the business plans were evaluated, AI's assessments were more aligned with the investor panel's average judgment than individual investors' assessments were—which suggests that AI can both generate strategic options at scale and evaluate them with a consistency exceeding that of human experts. Such findings don't mean AI is ready to replace your strategy team. But they do suggest that the cognitive work at the heart of strategy is no longer exclusively human territory.


In this article I'll lay out what this shift means in practice. First, I'll describe three ways AI can transform the strategy process, with examples from companies already using it to do so. Then I'll address the

skeptic's natural question—If every company has access to the same AI, how can it be a source of advantage?—and explain how firms can build durable competitive moats around strategy capabilities augmented by AI. Finally, I'll offer a practical playbook for leaders who want to start redesigning how their organizations make strategic decisions.

Three Ways AI Expands Strategic Thinking

To understand what changes with AI, it helps to think about what your organization has to do when making strategy. At its core, any strategic decision involves three cognitive tasks: *searching* for possible courses of action, *representing* the environment in which those actions will play out, and *aggregating* the judgments of the people involved in the decision. Let's look at each in turn.

From a handful of options to thousands. The most immediate constraint on any strategy process is in the search for options: How many alternatives can your team generate and evaluate? For most organizations the answer is surprisingly low. A typical strategic-planning cycle might surface a dozen ideas and then narrow the field to three or four, with team members spending most of their energy debating those finalists. The vast majority of the possibility space is never explored—not because it lacks promise but because human teams lack the time and capacity to take on that job.



Dimitris Ladopoulos combines art, math, and algorithms to create a study of forms and variations.

As my research suggests, the breakthrough with AI is that the cost of generating and screening options falls sharply, helping organizations explore far more of the possibility space. The shift is already visible in M&A. In one case, documented by McKinsey, a software company used a generative AI scouting system that combined semantic search with a database of more than 40 million public and private companies. Drawing on patents, filings, expert transcripts, and other structured

and unstructured data, the system surfaced and scored more than 500 acquisition targets in less than a day, narrowed the list to 15 serious leads, and supported three completed acquisitions within months.

The point here is not just speed. It is a qualitative shift in what “looking for options” means. Instead of asking a team to brainstorm plausible targets, a firm can now scan nearly the entirety of the relevant universe, using multiple strategic lenses simultaneously, and create a short list that no earlier process would have produced. The implication for leaders is straightforward: Before any major strategic choice—an acquisition, a new market entry, a product portfolio decision—you should task an AI system with generating a far larger set of alternatives than your team would normally consider, filtered by your specific criteria. Use your team’s judgment on the short list, not the long list. AI expands the search; humans choose.

From static frameworks to living models. The second constraint is representation: How does your organization model the environment it operates in? For most companies the answer is some combination of frameworks, spreadsheets, and quarterly reports—tools that are useful

but inherently static and simplified. They capture the most important dynamics but necessarily leave out enormous amounts of nuance, context, and real-time change.

AI makes it possible to work with far richer and more dynamic representations. MYbank, the digital lender within Ant Group, offers a striking example. Traditional banks evaluate small businesses' creditworthiness using just a handful of variables—revenue, collateral, credit history. Many small businesses lack one or more of these, leaving millions effectively “unscorable” and therefore unfundable. MYbank replaced that basic model with an AI system that draws on more than 3,000 variables, integrating transaction data, supply chain relationships, business network data, and even satellite imagery. The result is what the company calls its 3-1-0 model: three minutes to apply, one second for approval, zero human intervention. MYbank has extended credit to more than 53 million small businesses—72% of which were first-time borrowers—with a default rate around 1%.

The strategic insight here is about what happens when you replace a low-resolution model of your environment with a high-resolution one. Using its AI system, MYbank didn't just serve existing customers faster. It became able to see an entirely new customer segment—tens of millions of viable borrowers—that the old model had rendered invisible.

The same logic applies to any company whose models of demand, supply, or competitive dynamics were built for an earlier era. Consider Unilever's ice-cream business. Rather than relying on lagging quarterly forecasts, the company created an AI-driven model that integrates weather probabilities, demand signals, and telemetry from roughly 3 million connected freezers across 35 factories. The result is a continuously refreshed picture of where demand is forming and where supply may fall short—a living representation of the market that lets

managers intervene in real time rather than react after the fact. Unilever reports a 10% improvement in forecast accuracy in Sweden and sales increases of up to 30% in some regions from insights generated by AI-enabled freezers.

For leaders, the practical step is this: Identify the one or two strategic models your organization uses most heavily (your market map, your customer segmentation, your competitive positioning) and ask whether AI could make them richer, more current, and more granular. In most cases the answer will be yes. And the payoff won't just be better information. It will be the ability to see opportunities and risks that your current models are incapable of revealing.

From groupthink to structured challenge. The third constraint is aggregation: How does your organization combine the knowledge and judgment of multiple people into a decision? In theory, diverse teams should produce better decisions than any individual could alone. In practice, strategy meetings are shaped by hierarchy, politics, personality, and time limits. The most senior person's view carries disproportionate weight. Dissent can be risky. The group converges too quickly.

AI offers a different approach. Instead of depending on a roomful of people with uneven incentives to speak up, organizations can now orchestrate what might be called *synthetic deliberation*: AI-driven processes that pressure-test ideas without the social frictions of group dynamics.

The mechanism can be simple or sophisticated. McKinsey, for example, has described multiagent workflows in which a "creator" agent drafts an analysis and a "critic" agent systematically challenges it—a process that automates the rigorous back-and-forth that strong teams try to create

manually. At the more ambitious end, the BCG Henderson Institute has described LLM-powered strategic war games in which agents represent competitors, regulators, customers, and other stakeholders. Unlike in traditional rule-based simulations, these agents can generate less predictable, more humanlike responses, which forces management teams to confront scenarios they might never have imagined on their own.

The most rigorous evidence for AI's ability to improve collaborative ideation comes from a field experiment at Procter & Gamble. In a study involving 776 commercial and R&D professionals working on real innovation tasks, researchers found that individuals using AI matched the average quality of two-person teams working without it. Teams using AI were about 12% faster and were more likely to produce top-decile solutions. Perhaps most important, AI reduced the silo effect between commercial and R&D professionals, giving each function access to the other's perspective and helping both produce more-balanced ideas.

The lesson for leaders is to build structured AI challenges into your decision-making process. Assign AI agents to distinct roles: one to generate the strongest case for a strategic direction, one to argue against it, and one to simulate likely competitor responses. You can also create synthetic panels (customers, regional managers, regulators) to investigate how a decision might land with different stakeholders. The goal is to ensure that by the time your leadership team sits down to decide, the obvious weaknesses have been identified and the plan has been pressure-tested against a wider range of perspectives than any single meeting could accommodate.

None of this means AI outputs should be taken at face value. LLMs can produce confident-sounding analyses that are subtly wrong, internally

inconsistent, or built on fabricated evidence. Synthetic deliberation can surface useful challenges—but it can also generate plausible-seeming objections that miss the point entirely. The quality of AI-augmented strategy depends on the quality of the data fed into the process and, critically, the human judgment applied to its outputs. That isn't a reason to avoid these tools, but it is a reason to design the process carefully, and it's precisely why the role of the human strategist becomes more important, not less, in an AI-augmented world.

Building Advantage When Everyone Has AI

If you've read this far, you may be thinking: *But if my competitors can buy the same AI models that I have, how does any of this become a source of competitive advantage?* It's a fair question, and it has a clear answer if you look at it through the right lens.

The most useful analogy is the internet in 1995. By the mid-1990s, internet access was becoming widely available. The underlying technology was not proprietary. Yet the companies that built the most durable advantages of the next two decades—Amazon, Netflix, Google—did not win because they had a website. They won because they recognized that the internet wasn't something to just bolt onto existing operations; it was a new foundation on which to build business models, capabilities, and data assets that competitors could not easily replicate. The firms that treated the internet as a utility saw their offerings quickly become commoditized. The firms that treated it as a platform for reinvention created moats that lasted for decades.



Dimitris Ladopoulos

AI is at a similar inflection point. The models themselves are increasingly available. The advantage lies in what you build on top of them. Three approaches stand out.

Make AI smarter with your data. General-purpose AI models are trained on public information. They are, by definition, generic. The moment you fine-tune or ground a model with proprietary data—your customer histories, transaction patterns, or internal research—you create something that competitors cannot replicate simply by buying the same software. Morgan Stanley’s [AI assistant](#), which gives the company’s 16,000 financial advisers access to more than 100,000 internal research documents, reached 98% adoption and raised document-retrieval efficiency from 20% to 80%. Stripe’s fraud detection system, Radar, is another example at network scale. Trained on data from millions of businesses processing more than \$1.4 trillion in annual payments, there is a 92% chance that it has already seen any given card on its network. A smaller competitor using the same algorithms but lacking that data cannot match its accuracy.

Deploy AI in processes that only you have. Data is one moat. Process is another. When a company embeds AI into a proprietary workflow—one

that reflects its specific strategic logic, organizational knowledge, and competitive context—it creates a capability that’s difficult to observe from the outside and even harder to copy.

John Deere’s See & Spray system illustrates this concept well. The underlying technologies—computer vision and machine learning—are widely available. But Deere embedded them into precision-agriculture equipment that scans roughly 2,100 square feet per second, makes more than 5,000 spraying decisions per minute, and reduces herbicide use by 50% to 77%. The moat is not the algorithm. It’s the integration of AI into hardware, field-level data, and a dealer network far more expansive than that of rivals. What makes this approach strategic, not merely operational, is that See & Spray changes Deere’s competitive positioning: It transforms the company from an equipment manufacturer into a precision-agriculture platform, deepening customer lock-in and opening new revenue streams around data-driven farm management.

Move to the new frontier faster than your rivals. There is a well-known idea in strategy that operational effectiveness—doing the same things as other companies but better—is not the same as strategy. In stable times, that distinction holds. But in periods of technological discontinuity, the gap between firms that have reached the new performance frontier and firms that have not can be enormous.

Duolingo exemplifies the point. It launched AI-powered features—Roleplay, Explain My Answer, Video Call—soon after OpenAI’s GPT-4 became publicly available. But speed alone is not its moat. Duolingo’s real advantage is Birdbrain, a proprietary learning model trained on behavioral data from more than 500 million users completing roughly 1.25 billion exercises daily. The general-purpose model is the commodity. The learning system built around it, refined over years and

fed by behavioral data whose scale dwarfs that of competitors, is the moat.

AI can compress months of analysis into days, reveal opportunities that were invisible, and stress-test decisions that would otherwise go unchallenged. Firms that adopt those capabilities operate at a fundamentally different level of strategic quality. The advantage may be temporary in theory, but in practice it just keeps compounding: Better decisions lead to better positioning, which generates better data, which improves future decisions. The frontier will keep moving, so enduring advantage comes from the ability to reach it again and again, gaining strength each time.

The bottom line is this: As a technology, AI is becoming ubiquitous. But as a capability—embedded in your data, your processes, and your speed of execution—it is not. The companies that treat AI as a commodity to be purchased will find that it raises the bar for everyone without distinguishing them. The companies that treat it as a platform for building proprietary strategic capabilities will find that it is one of the most powerful sources of competitive advantage available today.

A Leader's Playbook for the AI Era

If you want to redesign how your organization makes its most important decisions, the playbook is straightforward.

Widen the option set before narrowing it. For every major strategic decision—an acquisition, a market entry, a product portfolio shift—use AI to generate a long list before the leadership team debates a short list. Most organizations narrow too early because exploring lots of options used to be expensive. It no longer is.

Replace static snapshots with living models. Pick the strategic representation your organization relies on most heavily—your customer segmentation, your competitive map, your demand forecast—and upgrade it from a periodic document to a dynamic model that updates as conditions change. The goal is to create a map rich enough to reveal what your current framework structurally cannot.

Institutionalize structured challenge. Good strategy depends not just on generating ideas but also on rigorously testing them. In most organizations that testing is still shaped by hierarchy, diplomacy, and the human tendency to converge too quickly. Don't leave dissent to chance. Design it into the process.

Build creator-critic-competitor workflows into every major commitment: Assign one AI agent to make the strongest case for a plan, one to dismantle it, and one to simulate how rivals or regulators would respond. Make structured challenge as routine as financial modeling.

Redefine the strategist's role from analyst to architect. In a world of tight cognitive bounds, strategists spend most of their time gathering data, building spreadsheets, and populating simplified frameworks. As AI starts to unbound cognitive constraints, that baseline cognitive work can increasingly be automated—and the key resource shifts from analysis to imagination, from populating frameworks to designing the process itself.

All of those shifts demand a new talent profile: the hybrid strategist. Leaders who fit this profile know how to frame a strategic question, can design an AI workflow to answer it, and, above all, understand where the machine must yield to human judgment.

To become a hybrid strategist, you'll need to change what you demand in the room where decisions are made. Before any major strategic proposal reaches the leadership table, require three things: the AI-expanded list of alternatives that were considered and rejected, a current model of the competitive environment, and the results of a structured critique. When the executive committee raises the standard for what a credible recommendation looks like, the rest of the organization will follow.

Finally, make AI use normal. Morgan Stanley's AI assistant reached 98% adoption among 16,000 financial advisers because it was woven into the proprietary research workflows that they already depended on. The tool made their daily work clearly better, and adoption took care of itself. Start with the same logic: Find the decision process where your people already feel the pain, rebuild it with AI, and make the new workflow the default. Culture usually follows usefulness.

...

Picture your next strategy offsite. But this time, imagine it unbounded.

Your team arrives having already reviewed an AI-generated scan of hundreds of possibilities, filtered to the most promising. The market model on the wall is current, not six months stale. Before anyone presents a recommendation, it has survived a structured critique—devil's advocate, simulated competitors, skeptical customers. The conversation in the room is sharper, the options are stronger, and the time that used to go to basic analysis now goes to the questions only humans can answer: What do we believe? What are we willing to risk? What kind of company do we want to become?

A version of this article appeared in the [September–October 2026](#) issue of Harvard Business Review.



Felipe A. Csaszar is the Alexander M. Nick Professor and the chair of the strategy area at the University of Michigan's Ross School of Business.